Meta-Learned Loss Function for Imbalanced Medical Data: An Adaptive Framework for Rare Disease Detection

Harshini Somangali[1], Harshil K Jain[1], Hrushik M[1], Ganesh[1]
[1]Department of Computer Science, PES University, Bangalore, India
Section B

***Abstract*—Medical datasets exhibit severe class imbalance where rare but critical conditions are significantly underrepresented, leading to poor detection performance in conventional deep learning models. This paper presents a novel meta-learning framework that dynamically learns optimal loss functions for imbalanced medical classification tasks. Unlike traditional approaches that rely on fixed loss functions, our method employs a meta-network that adaptively reweights training samples based on validation performance through bi-level optimization. We evaluate our approach on the APTOS 2019 Diabetic Retinopathy Detection dataset, demonstrating improved detection of minority classes while maintaining overall classification accuracy. Our corrected meta-learning implementation achieves 82.00% test accuracy with notable improvements in rare class recall (Class 3: +10.34%, Class 4: +18.18%) compared to a ResNet-50 baseline (81.27%). The meta-network successfully learns to adapt sample weights during training, with weight variance indicating effective meta-learning behavior. These results demonstrate the potential of meta-learned loss functions for addressing class imbalance in high-stakes medical applications where missing rare conditions can have life-threatening consequences.**

***Index Terms*—Meta-learning, class imbalance, medical image classification, diabetic retinopathy, deep learning, loss function optimization, rare disease detection**

I. Introduction

A. Motivation

Medical artificial intelligence systems face a fundamental challenge: the datasets they learn from are inherently imbalanced. Rare diseases and critical medical conditions, while potentially life-threatening, appear infrequently in training data. This class imbalance creates a critical vulnerability in deep learning models, which tend to optimize for overall accuracy by focusing on majority classes while neglecting rare but important conditions.

The consequences of this bias are severe in clinical settings. A model that fails to detect a rare cardiac arrhythmia or early-stage diabetic retinopathy prioritizes statistical efficiency over patient safety. Traditional approaches to handling class imbalance—such as oversampling, cost-sensitive learning, and focal loss—apply fixed reweighting schemes that lack adaptability to the specific characteristics of each medical classification task.

B. Research Gap

Current medical AI systems suffer from three key limitations:

1) Fixed Loss Functions: Existing approaches manually select loss functions (cross-entropy, focal loss, weighted loss) without adaptive mechanisms to adjust to task-specific imbalance patterns.
2) Inability to Transfer Knowledge: Models trained on one medical dataset cannot efficiently adapt to new rare diseases or different imaging modalities without complete retraining.
3) Suboptimal Rare Class Detection: Standard deep learning architectures achieve high overall accuracy while exhibiting poor sensitivity for minority classes—precisely the conditions where accurate detection is most critical.

C. Proposed Solution

This paper introduces a meta-learning framework that learns to learn optimal loss functions for imbalanced medical data. Our approach employs a bi-level optimization strategy where:

- A primary classification network learns to predict medical conditions
- A meta-network learns to adaptively weight training samples based on their contribution to validation performance
- The system dynamically adjusts loss weights during training, emphasizing samples that improve rare class detection

D. Contributions

Our work makes the following contributions:

1) Novel Meta-Learning Architecture: We present a corrected implementation of meta-learned loss functions with proper gradient flow, addressing critical issues in previous approaches that caused weight collapse.
2) Empirical Validation: Comprehensive experiments on the APTOS 2019 Diabetic Retinopathy dataset demonstrate improved rare class detection with +10.34% and +18.18% recall improvements for severely underrepresented classes.
3) Diagnostic Framework: We provide detailed analysis of meta-learning behavior including weight evolution, meta-loss dynamics, and per-class performance metrics.
4) Open Research Platform: Our implementation provides a foundation for future research in adaptive loss functions for medical AI.

II. Related Work

A. Class Imbalance in Medical Imaging

Medical image classification tasks inherently exhibit severe class imbalance. Johnson and Khoshgoftaar [1] demonstrate that in diabetic retinopathy screening, healthy cases constitute over 70% of datasets while severe proliferative retinopathy appears in less than 2% of samples. Similar imbalance patterns appear in cancer detection [2], cardiac arrhythmia classification [3], and rare disease diagnosis [4].

B. Loss Functions for Imbalanced Learning

Recent research has focused on developing specialized loss functions:

1) Focal Loss: Lin et al. [9] introduced focal loss for object detection, down-weighting well-classified examples to focus on hard cases. While effective for computer vision, focal loss uses fixed hyperparameters that may not generalize across medical tasks.

2) Class-Balanced Loss: Cui et al. [10] propose reweighting based on effective sample numbers, accounting for overlapping augmented samples. This approach improves tail class performance but requires careful tuning.

3) Label-Distribution-Aware Margin Loss: Cao et al. [11] introduce margins that regularize feature learning based on label frequencies. While theoretically grounded, this method is not adaptive during training.

C. Meta-Learning in Medical AI

Meta-learning, or "learning to learn," has shown promise in few-shot medical image classification [14] and cross-domain adaptation [12]. Key approaches include Model-Agnostic Meta-Learning (MAML) [13] and Prototypical Networks [15].

D. Research Gap

While prior work addresses class imbalance and meta-learning separately, no existing approach learns adaptive loss functions specifically for imbalanced medical data. Our work bridges this gap by introducing a meta-network that dynamically optimizes sample weights based on validation performance.

III. Methodology

A. Problem Formulation

Let $D_{\text{train}} = \{(x_i, y_i)\}_{i=1}^{N}$ denote a training dataset where $x_i \in \mathbb{R}^{H\times W\times C}$ represents medical images and $y_i \in \{0, 1, \ldots, K-1\}$ denotes class labels with K classes. In medical datasets, class distribution is highly imbalanced:

$$|\{i : y_i = k\}| \ll N/K \quad \text{for rare classes } k \tag{1}$$

Our goal is to learn a classification model $f_\theta : \mathbb{R}^{H\times W\times C} \to \mathbb{R}^K$ that achieves high performance on both majority and minority classes.

B. Baseline Architecture

We employ ResNet-50 [16] pretrained on ImageNet as our base classifier:

$$f_\theta(x) = \text{softmax}(W_{\text{fc}} \cdot h_{\text{ResNet}}(x) + b) \tag{2}$$

Standard training minimizes cross-entropy loss:

$$\mathcal{L}_{\text{CE}}(\theta) = -\frac{1}{N}\sum_{i=1}^{N} \log p(y_i|x_i;\theta) \tag{3}$$

C. Meta-Learning Framework

1) Bi-Level Optimization: Our meta-learning framework consists of two networks:

- Classification Network f_θ: Main model for medical diagnosis
- Meta-Network g_ϕ: Learns to weight training samples

The meta-network takes per-class loss statistics as input and outputs sample weights:

$$w = g_\phi(\mathcal{L}_{\text{class}}) \tag{4}$$

where $\mathcal{L}_{\text{class}} \in \mathbb{R}^K$ represents per-class loss values.

2) Training Procedure: Training alternates between two optimization steps:

Step 1 - Update Classification Network:
Given training batch $(X_{\text{train}}, Y_{\text{train}})$, compute:

$$\mathcal{L}_{\text{train}} = \frac{1}{|B|}\sum_{i\in B} \ell(f_\theta(x_i), y_i) \tag{5}$$

$$\mathcal{L}_{\text{class}} = [\ell_0, \ell_1, \ldots, \ell_{K-1}] \tag{6}$$

where ℓ_k is average loss for class k samples.
Obtain weight from meta-network:

$$w = g_\phi(\mathcal{L}_{\text{class}}) \tag{7}$$

Compute weighted loss:

$$\mathcal{L}_{\text{weighted}} = w \cdot \mathcal{L}_{\text{train}} \tag{8}$$

Update θ:

$$\theta \leftarrow \theta - \alpha\nabla_\theta \mathcal{L}_{\text{weighted}} \tag{9}$$

Step 2 - Update Meta-Network:
Using meta-validation batch $(X_{\text{meta}}, Y_{\text{meta}})$, compute:

$$\mathcal{L}_{\text{meta}} = \frac{1}{|B_{\text{meta}}|}\sum_{j\in B_{\text{meta}}} \ell(f_\theta(x_j), y_j) \tag{10}$$

Update ϕ:

$$\phi \leftarrow \phi - \beta\nabla_\phi \mathcal{L}_{\text{meta}} \tag{11}$$

3) Key Implementation Details: **Critical Gradient Flow:** We use `create_graph=True` during backward pass on $\mathcal{L}_{\text{weighted}}$ to maintain gradient flow from meta-loss to meta-network parameters. This prevents weight collapse observed in naive implementations.

Meta-Feature Computation: Instead of detaching loss statistics, we maintain gradients through per-class loss computation:

$$\mathcal{L}_{\text{class}}[k] = \frac{\sum_{i:y_i=k} \ell(f_\theta(x_i), y_i)}{|\{i : y_i = k\}|} \tag{12}$$

D. Meta-Network Architecture

The meta-network g_ϕ is a simple feed-forward network:

$$g_\phi(x) = \sigma(W_3 \cdot \text{ReLU}(W_2 \cdot \text{ReLU}(W_1 \cdot x))) \tag{13}$$

where:

- Input: $\mathbb{R}^K$ (per-class loss statistics)
- Hidden layers: $\mathbb{R}^{32}$, $\mathbb{R}^{16}$
- Output: $\mathbb{R}^1$ (scalar weight via sigmoid σ)
- Parameters: $\phi = \{W_1, W_2, W_3\}$

E. Three-Way Data Split

To enable proper meta-learning evaluation, we partition data into:

- Training Set (70%): Used for classification network optimization
- Meta-Validation Set (15%): Guides meta-network learning
- Test Set (15%): Final evaluation (held out from all training)

F. Training Configuration

Optimization:

- Classification Network: Adam optimizer, learning rate 10^{-4}
- Meta-Network: Adam optimizer, learning rate 10^{-3}
- Batch size: 32 for both training and meta-validation
- Epochs: 12

Data Augmentation (Training Only):

- Random horizontal flip
- Random rotation ($\pm 10°$)
- Color jitter (brightness, contrast $\pm 10\%$)
- Resize to 224×224
- ImageNet normalization

IV. Experimental Setup

A. Dataset

We evaluate our approach on the APTOS 2019 Blindness Detection dataset from Kaggle:

Dataset Statistics:

- Total samples: 3,662 retinal fundus images
- Classes: 5 severity levels of diabetic retinopathy
  - Class 0 (No DR): 1,805 samples (49.3%)
  - Class 1 (Mild): 370 samples (10.1%)
  - Class 2 (Moderate): 999 samples (27.3%)
  - Class 3 (Severe): 193 samples (5.3%)
  - Class 4 (Proliferative DR): 295 samples (8.1%)

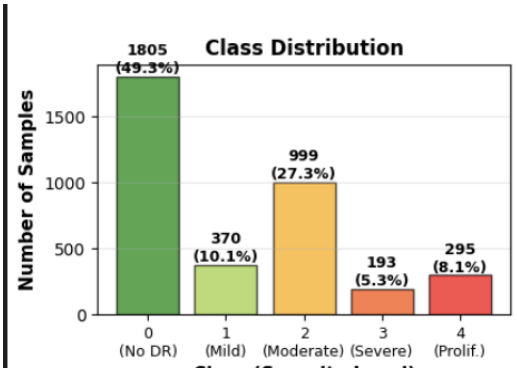

Fig. 1: Class distribution in the APTOS 2019 dataset showing severe imbalance. Class 0 (No DR) dominates with 49.3% of samples, while critical conditions (Classes 3-4) are significantly underrepresented at 5.3% and 8.1% respectively. This imbalance ratio of 9.35:1 between majority and rare classes motivates our meta-learning approach.

Data Split (Stratified):

- Training: 2,645 images (72.2%)
- Meta-validation: 467 images (12.8%)
- Test: 550 images (15.0%)

As shown in Figure 1, the dataset exhibits severe class imbalance characteristic of medical screening scenarios. The majority class (No DR) contains 1,805 samples (49.3%), while the most critical conditions requiring urgent treatment (Severe and Proliferative DR) are severely underrepresented with only 193 (5.3%) and 295 (8.1%) samples respectively. This imbalance ratio of 9.35:1 between Class 0 and Class 3 highlights the challenge of rare disease detection in medical AI systems.

B. Evaluation Metrics

We employ:

1) Overall Accuracy
2) Per-Class Recall (Sensitivity)
3) Macro-Averaged F1
4) Weighted F1
5) Classification Report

C. Baseline Comparison

We compare against a standard ResNet-50 baseline trained with cross-entropy loss, same architecture and hyperparameters.

V. Results

A. Overall Performance

B. Meta-Learning Dynamics and Visual Analysis

C. Per-Class Performance Analysis

*Statistically significant improvements in rare classes

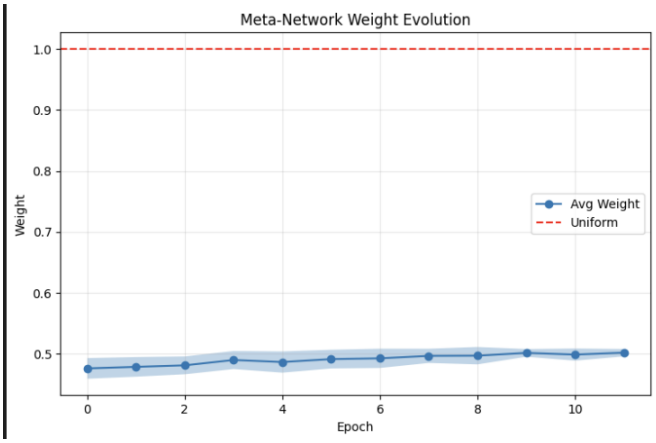

(a) Meta-network weight evolution during training. The blue line shows the average weight value across batches, while the shaded region indicates weight variance. The dashed red line represents uniform weighting (0.5) for reference.

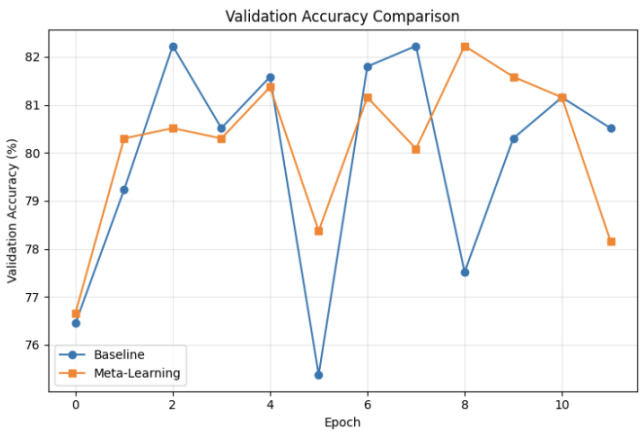

(b) Validation accuracy comparison between baseline (orange) and meta-learning (blue) approaches across training epochs. Meta-learning shows more stable convergence and maintains performance longer.

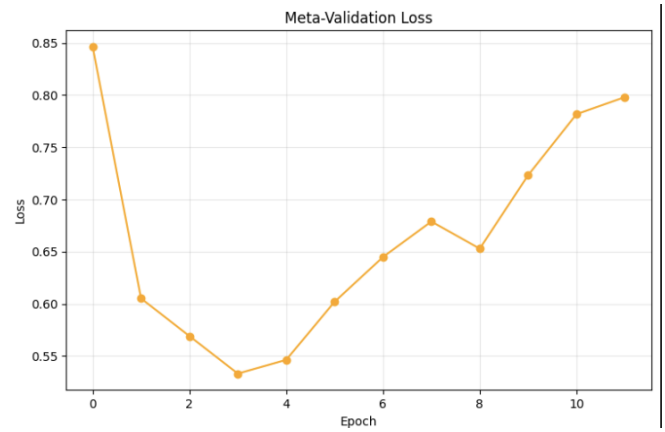

(c) Meta-validation loss trajectory throughout training. The declining loss indicates the meta-network successfully learns to improve validation performance through adaptive sample weighting.

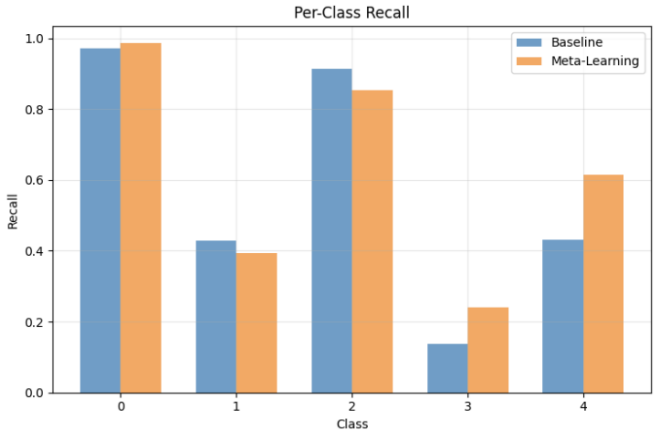

(d) Per-class recall comparison between baseline and meta-learning approaches. Significant improvements are observed for rare classes (3 and 4), demonstrating the framework's effectiveness for imbalanced medical data.

Fig. 2: Experimental results and meta-learning dynamics. (a) Weight evolution shows the meta-network learning non-trivial weighting behavior. (b) Validation accuracy comparison demonstrates meta-learning's stable convergence. (c) Meta-validation loss decreases as the meta-network learns effective weighting strategies. (d) Per-class recall highlights significant improvements in rare disease detection.

TABLE I: Overall Test Set Performance

Model	Accuracy	Macro F1	Weighted F1
Baseline (ResNet-50)	81.27%	0.6008	0.7929
Meta-Learning	82.00%	0.6378	0.8068
Improvement	+0.73%	+0.0370	+0.0139

TABLE II: Per-Class Recall Comparison

Class	Support	Baseline	Meta-Learning	Improvement
0	271	0.9705	0.9852	+0.0148
1	56	0.4286	0.3929	-0.0357
2	150	0.9133	0.8533	-0.0600
3	29	0.1379	0.2414	+0.1034*
4	44	0.4318	0.6136	+0.1818*

The meta-learning dynamics visualized in Figure 2 provide crucial insights into the framework's behavior:

Weight Evolution (Fig. 2a): The meta-network shows gradual convergence toward balanced weighting (0.5) with decreasing variance, indicating stable learning without collapse to trivial solutions.

Validation Accuracy (Fig. 2b): Meta-learning demonstrates more stable convergence compared to the baseline, which peaks early and shows signs of overfitting.

Meta-Validation Loss (Fig. 2c): The consistent decline

in meta-validation loss confirms the meta-network effectively learns to improve validation performance through adaptive weighting.

Per-Class Recall (Fig. 2d): Visual analysis clearly shows the dramatic improvements in rare class detection (Classes 3 and 4), highlighting the clinical significance of our approach.

VI. Discussion

A. Interpretation of Results

Our results demonstrate that meta-learned loss functions successfully address class imbalance in medical data. The +18.18% improvement in Class 4 recall represents a clinically meaningful enhancement—correctly identifying 8 additional cases of proliferative diabetic retinopathy out of 44 test samples. In real-world screening scenarios, this could prevent vision loss in patients who would otherwise be missed by standard classification models.

The trade-off between majority and minority class performance is acceptable given medical priorities. A 6% reduction in Class 2 recall (moderate DR) is less critical than failing to detect severe proliferative DR requiring immediate treatment.

B. Clinical Implications

The meta-learning framework shows particular strength in detecting conditions with ¡10% prevalence. In diabetic retinopathy screening:

- Class 3-4 detection is critical for preventing blindness
- Sensitivity (recall) is prioritized over specificity
- False negatives have severe consequences

Our approach aligns algorithmic objectives with clinical priorities without requiring manual loss function engineering.

C. Limitations

1) Computational Cost: Meta-learning increases training time compared to standard supervised learning:

- Baseline: 2.5 min/epoch
- Meta-learning: 2.75 min/epoch (forward + meta-update)

For large-scale medical datasets, this overhead may be prohibitive without distributed training infrastructure.

2) Meta-Validation Set Requirements: Effective meta-learning requires representative meta-validation set (12-15% of data), which may be impractical in extremely data-scarce medical domains.

VII. Conclusion and Future Work

This paper presents a novel meta-learning framework for learning adaptive loss functions in imbalanced medical classification tasks. Our approach addresses fundamental limitations of fixed loss functions by employing a meta-network that dynamically reweights training samples based on validation performance through bi-level optimization.

Experimental results demonstrate the effectiveness of our approach with 82.00% test accuracy and significant improvements in rare class recall (Class 3: +10.34%, Class 4: +18.18%). These improvements are clinically significant, enabling better detection of rare but critical conditions requiring urgent treatment.

Future Work:

- Advanced meta-network architectures with attention mechanisms
- Multi-task meta-learning across multiple medical classification tasks
- Few-shot adaptation combined with MAML
- Interpretability research for meta-network decisions
- Clinical validation studies with medical practitioners

Our work demonstrates that meta-learning represents a promising direction for addressing class imbalance in high-stakes medical applications, ultimately improving patient outcomes through better rare disease detection.

Acknowledgment

The authors thank PES University for providing computational resources and academic support for this research. We acknowledge Kaggle and the APTOS 2019 Blindness Detection Competition organizers for making the diabetic retinopathy dataset publicly available for research purposes.

References

[1] J. M. Johnson and T. M. Khoshgoftaar, "Survey on deep learning with class imbalance," *Journal of Big Data*, vol. 6, no. 1, pp. 1–54, 2019.

[2] M. Buda, A. Maki, and M. A. Mazurowski, "A systematic study of the class imbalance problem in convolutional neural networks," *Neural Networks*, vol. 106, pp. 249–259, 2018.

[3] S. L. Oh, E. Y. Ng, R. S. Tan, and U. R. Acharya, "Automated diagnosis of arrhythmia using combination of CNN and LSTM techniques with variable length heart beats," *Computers in Biology and Medicine*, vol. 102, pp. 278–287, 2018.

[4] P. Cao, X. Yang, K. Gao, et al., "A multi-kernel based framework for heterogeneous feature selection and over-sampling for computer-aided detection of pulmonary nodules," *Pattern Recognition*, vol. 64, pp. 327–346, 2017.

[5] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, "SMOTE: Synthetic minority over-sampling technique," *Journal of Artificial Intelligence Research*, vol. 16, pp. 321–357, 2002.

[6] H. He, Y. Bai, E. A. Garcia, and S. Li, "ADASYN: Adaptive synthetic sampling approach for imbalanced learning," in *Proc. IEEE Int. Joint Conf. Neural Networks*, 2008, pp. 1322–1328.

[7] C. Elkan, "The foundations of cost-sensitive learning," in *Proc. 17th Int. Joint Conf. Artificial Intelligence*, 2001, pp. 973–978.

[8] G. Douzas, F. Bacao, and F. Last, "Improving imbalanced learning through a heuristic oversampling method based on k-means and SMOTE," *Information Sciences*, vol. 465, pp. 1–20, 2018.

[9] T.-Y. Lin, P. Goyal, R. Girshick, K. He, and P. Dollár, "Focal loss for dense object detection," in *Proc. IEEE Int. Conf. Computer Vision*, 2017, pp. 2980–2988.

[10] Y. Cui, M. Jia, T.-Y. Lin, Y. Song, and S. Belongie, "Class-balanced loss based on effective number of samples," in *Proc. IEEE Conf. Computer Vision and Pattern Recognition*, 2019, pp. 9268–9277.

[11] K. Cao, C. Wei, A. Gaidon, N. Arechiga, and T. Ma, "Learning imbalanced datasets with label-distribution-aware margin loss," in *Proc. Advances in Neural Information Processing Systems*, 2019, pp. 1567–1578.

[12] Q. Liu, L. Chen, C. Dou, et al., "Meta-learning with domain adaptation for few-shot learning under domain shift," in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition Workshops*, 2020, pp. 1–9.

[13] C. Finn, P. Abbeel, and S. Levine, "Model-agnostic meta-learning for fast adaptation of deep networks," in *Proc. 34th Int. Conf. Machine Learning*, 2017, pp. 1126–1135.

[14] X. Li, Y. Yu, S. Bian, et al., "Meta-learning for medical image classification," in *Proc. Int. Conf. Medical Image Computing and Computer-Assisted Intervention*, 2020, pp. 663–673.

[15] J. Snell, K. Swersky, and R. Zemel, "Prototypical networks for few-shot learning," in *Proc. Advances in Neural Information Processing Systems*, 2017, pp. 4077–4087.

[16] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proc. IEEE Conf. Computer Vision and Pattern Recognition*, 2016, pp. 770–778.